Lecture 15:
Solving and Estimating STATIC Games of Incomplete Information

2021 Econometric Society Summer School in
Dynamic Structural Econometrics

Fedor Iskhakov, Australian National University
John Rust, Georgetown University
Bertel Schjerning, University of Copenhagen

University of Bonn
16-22 August, 2021

Road Map for Lectures on Games

Lecture 15: Estimation of static games with multiple equilibria

- Methods: NFXP, MPEC, CCP and NPL
- Example: Simple static entry model
- Explicit focus: Multiple Equilibira

Next: Solving and estimating directional dynamic games

- Example: Dynamic model of Bertrand duopoly competition and cost reducing investments
- Huge multiplicity of Equilibria
- Full solution method: Recursive Lexicographic Search (RLS)
- Estimation method: MLE using NRLS (implemented using Branch and Bounds algorithm)
- Compare with: MPEC, CCP estimator and Nested Pseudo Likelihood

Estimating Discrete-Choice Games of Incomplete Information

Estimating Discrete-Choice Games of Incomplete Information

- Aguirregabiria and Mira (2007): NPL (Recursive 2-Step)
- Bajari, Benkard and Levin (2007): 2-Step Minimum Distance Estimator
- Pakes, Ostrovsky and Berry (2007): Various 2-Step (PML, MoM, min χ^2)
- Pesendorfer and Schmidt-Dengler (2008): 2-Step Least Squares
- Pesendorfer and Schmidt-Dengler (2010): comments on AM (2007)
- Kasahara and Shimotsu (2012): Modified NPL
- Su (2013), Egesdal, Lai and Su (2015): Constrained Optimization

Example: Static Game Entry of Incomplete Information

- Two firms: a and b
- Actions: each firm has two possible actions:

$$d_a = \begin{cases} 1, & \text{if firm } a \text{ choose to enter the market} \\ 0, & \text{if firm } a \text{ choose not to enter the market} \end{cases} \tag{1}$$

$$d_b = \begin{cases} 1, & \text{if firm } b \text{ choose to enter the market} \\ 0, & \text{if firm } b \text{ choose not to enter the market} \end{cases} \tag{2}$$

Example: Static Entry Game of Incomplete Information

Utility: Ex-post payoff to firms

$$u_a(d_a, d_b, x_a, \epsilon_a) = \begin{cases} [\alpha + d_b * (\beta - \alpha)]x_a + \epsilon_{a1}, & \text{if } d_a = 1 \\ 0 + \epsilon_{a0}, & \text{if } d_a = 0 \end{cases}$$

$$u_b(d_a, d_b, x_a, \epsilon_b) = \begin{cases} [\alpha + d_a * (\beta - \alpha)]x_b + \epsilon_{b1}, & \text{if } d_b = 1 \\ 0 + \epsilon_{b0}, & \text{if } d_b = 0 \end{cases}$$

- (α, β): structural parameters to be estimated
- (x_a, x_b): firms' observed types; **common knowledge**
- $\epsilon_a = (\epsilon_{a0}, \epsilon_{a1}), \epsilon_b = (\epsilon_{b0}, \epsilon_{b1})$: firms' unobserved types, **private information**
- (ϵ_a, ϵ_b) are observed only by each firm, but not by their opponent firm nor by the econometrician

Example: Static Entry Game of Incomplete Information

- Assume the error terms (ϵ_a, ϵ_b) have a standardized type III extreme value distribution
- A Bayesian Nash equilibrium (p_a, p_b) satisfies

$$
\begin{aligned}
p_a &= \frac{\exp[p_b \beta x_a + (1 - p_b)\alpha x_a]}{1 + \exp[p_b \beta x_a + (1 - p_b)\alpha x_a]} \\
&= \frac{1}{1 + \exp[-\alpha x_a + p_b x_a(\alpha - \beta)]} \\
&\equiv \Psi_a(p_b, x_a; \alpha, \beta)
\end{aligned}
$$

$$
\begin{aligned}
p_b &= \frac{1}{1 + \exp[-\alpha x_b + p_a x_b(\alpha - \beta)]} \\
&\equiv \Psi_b(p_a, x_b; \alpha, \beta)
\end{aligned}
$$

Static Game Example: Parameters

We consider a very contestable game throughout

- Monopoly profits: $\alpha * x_j = 5 * x_j$
- Duopoly profits: $\beta * x_j = -11 * x_j$
- Firm types: $(x_a, x_b) = (0.52, 0.22)$

Static Game Example: Three Bayesian Nash Equilibria

Figure: Equilibria at intersections of best response functions

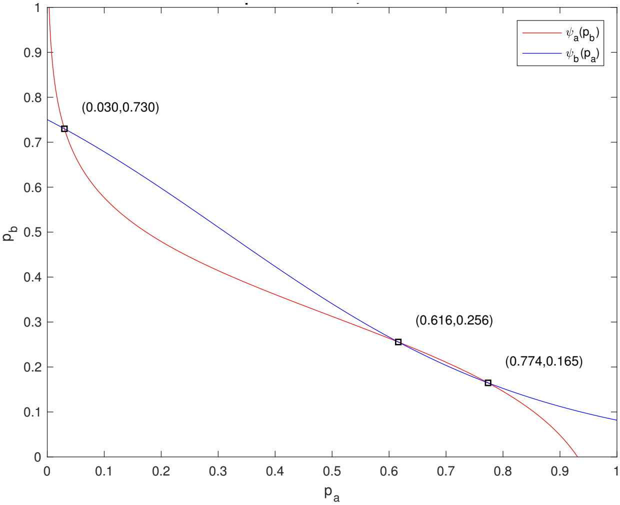

Static Game Example: Solving for Equilibria

Figure: Fixed points on second order best response function

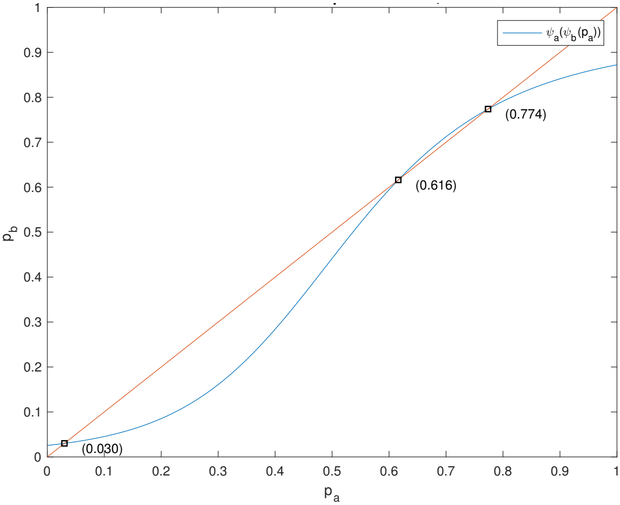

Static Game Example: Solving for Equilibria

Solution method: Combination of succesive approximations and bisection algorithm

Succesive approximations (SA)

- Converge to nearest stable equilibrium.
- Start SA at $p_a = 0$ and $p_a = 1$.
- Unique equilibrium (K=1): SA will converge to it from anywhere.
- Three equilibria (K=3): Two will be stable, and one will be unstable.
- More equilibria (K>3): Not in this model.

Bisection method

- Use this to find the unstable equilibrium (if K=3).
- The bisection method that repeatedly bisects an interval and then selects a subinterval in which the fixed point (or root) must lie.
- The two stable equilibria, defines the initial interval to search over.
- The bisection method is a very simple and robust method, but it is also relatively slow.

Static Game Example: Solving for Equilibria

Figure: Bisection method

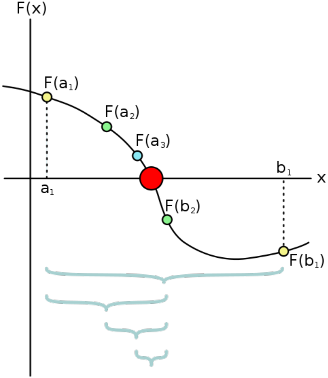

A few steps of the bisection method applied over the starting range [a1;b1]. The bigger red dot is the root of the function.

Static Game Example: Data Generation and Identification

- Data Generating Process (DGP): the data are generated by a single equilibrium
- The two players use the same equilibrium to play 1000 times
- Data: $X = \{(d_a^i, d_b^i)_{i=1}^{1000}, (x_a, x_b) = (0.52, 0.22)\}$
- Given data X, we want to recover structural parameters α and β

Static Game Example: Maximum Likelihood Estimation

- Maximize the likelihood function

$$
\begin{aligned}
\max_{\alpha,\beta} \quad \log \quad & \mathcal{L}(p_a(\alpha,\beta);X) \\
= \quad & \sum_{i=1}^{N} \left(d_a^i * \log(p_a(\alpha,\beta)) + (1-d_a^i*)\log(1-p_a(\alpha,\beta))\right) \\
+ \quad & \sum_{i=1}^{N} \left(d_b^i * \log(p_b(\alpha,\beta)) + (1-d_b^i*)\log(1-p_b(\alpha,\beta))\right)
\end{aligned}
$$

- $p_a(\alpha,\beta)$ and $p_a(\alpha,\beta)$ are the solutions of the Bayesian-Nash Equilibrium equations

$$
\begin{aligned}
p_a &= \frac{1}{1+\exp[-\alpha x_a + p_b x_a(\alpha-\beta)]} \equiv \Psi_a(p_b, x_a; \alpha, \beta) \\
p_b &= \frac{1}{1+\exp[-\alpha x_b + p_a x_b(\alpha-\beta)]} \equiv \Psi_b(p_a, x_b; \alpha, \beta)
\end{aligned}
$$

Static Game Example: MLE via NFXP

- Outer Loop
  - Choose (α, β) to maximize the likelihood function

$$\log \mathcal{L}(p_a(\alpha, \beta), p_b(\alpha, \beta); X)$$

- Inner loop:
  - For a given (α, β), solve the BNE equations for **ALL** equilibria: $(p_a^k(\alpha, \beta), p_b^k(\alpha, \beta)), k = 1, ..., K$
  - Choose the equilibrium that gives the highest likelihood value:

$$k^* = \arg \max_{k=1,\dots,K} \log \mathcal{L}(p_a^k(\alpha, \beta), p_b^k(\alpha, \beta); X)$$

such that

$$(p_a(\alpha, \beta), p_b(\alpha, \beta)) = (p_a^{k*}(\alpha, \beta), p_b^{k*}(\alpha, \beta))$$

NFXP's Likelihood as a Function of (α, β) - Eq 1

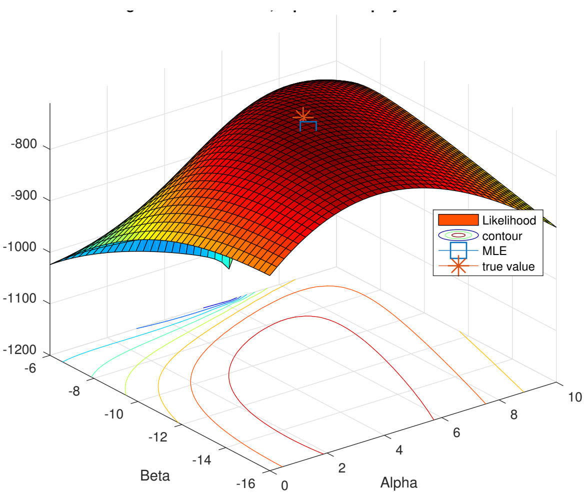

Figure: Data generated from equilibrium 1

NFXP's Likelihood as a Function of (α, β) - Eq 2

Figure: Data generated from equilibrium 2

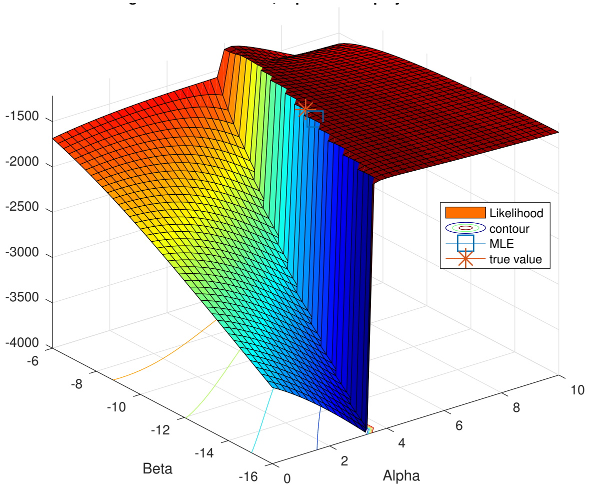

NFXP's Likelihood as a Function of (α, β) - Eq 3

Figure: Data generated from equilibrium 3

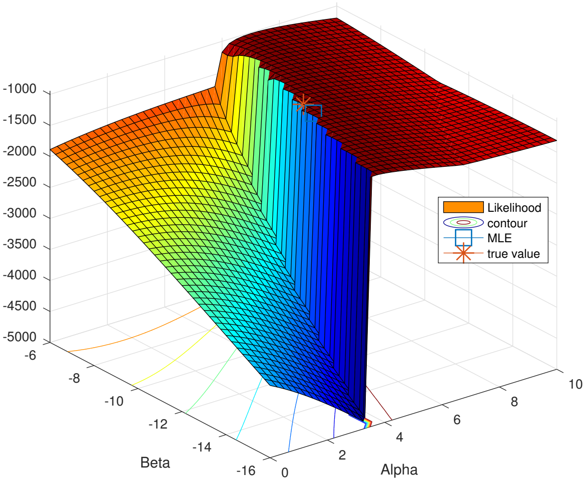

Monte Carlo Results: NFXP with Eq1

Figure: Data generated from equilibrium 1

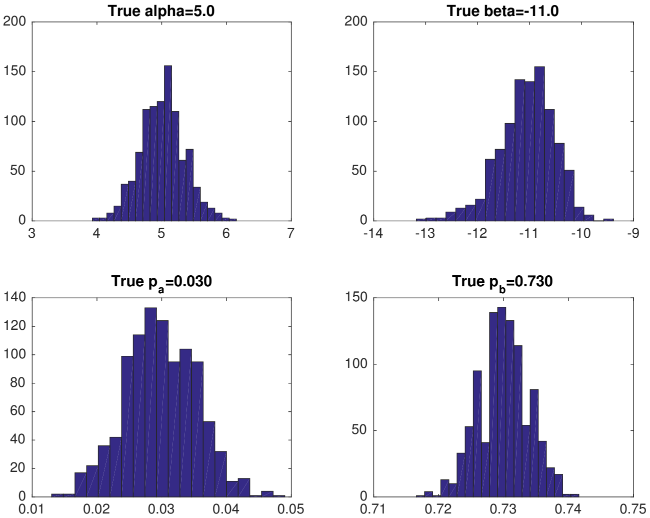

Monte Carlo Results: NFXP with Eq2

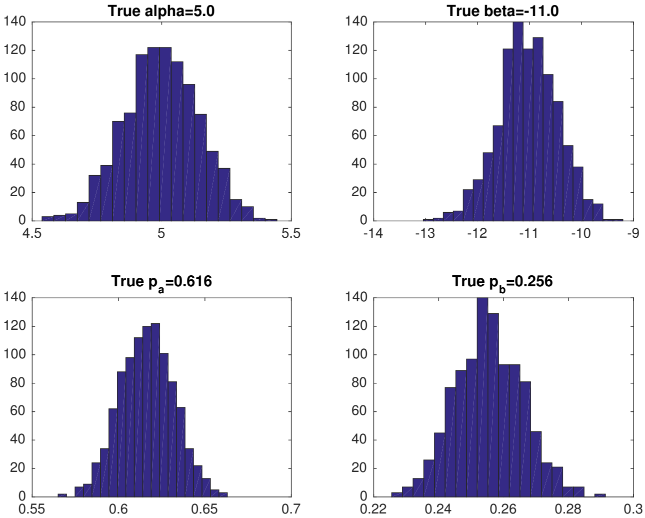

Figure: Data generated from equilibrium 2

Monte Carlo Results: NFXP with Eq3

Figure: Data generated from equilibrium 3

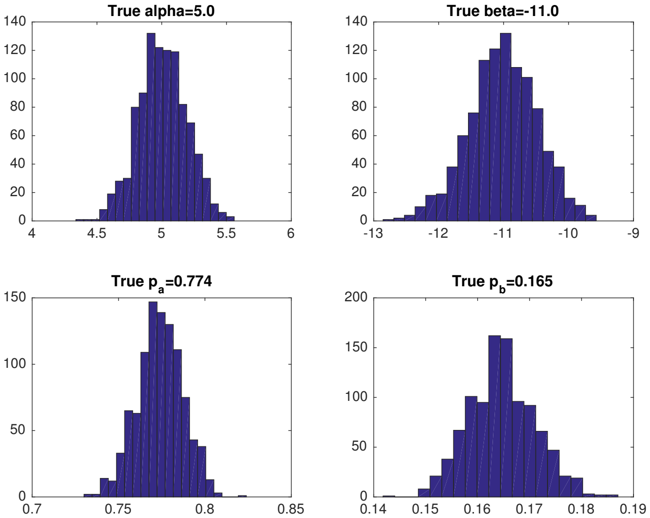

Constrained Optimization Formulation for Maximum Likelihood Estimation

- Maximize the likelihood function

$$
\begin{aligned}
\max_{\alpha,\beta,p_a,p_b} \quad \log \quad & \mathcal{L}(p_a; X) \\
= \quad & \sum_{i=1}^{N} \left(d_a^i * \log(p_a) + (1 - d_a^i *) \log(1 - p_a) \right) \\
+ \quad & \sum_{i=1}^{N} \left(d_b^i * \log(p_b) + (1 - d_b^i *) \log(1 - p_b) \right)
\end{aligned}
$$

- Subject to p_a and p_a are the solutions of the Bayesian-Nash Equilibrium equations

$$
\begin{aligned}
p_a &= \frac{1}{1 + \exp[-\alpha x_a + p_b x_a(\alpha - \beta)]} \\
p_b &= \frac{1}{1 + \exp[-\alpha x_b + p_a x_b(\alpha - \beta)]} \\
0 &\leq p_a, p_b \leq 1
\end{aligned}
$$

Monte Carlo Results: MPEC with Eq1

Figure: Data generated from equilibrium 1

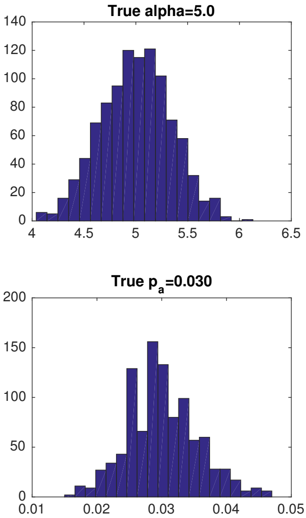

Monte Carlo Results: MPEC with Eq2

Figure: Data generated from equilibrium 2

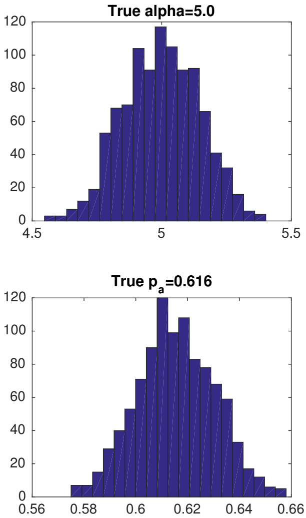

Monte Carlo Results: MPEC with Eq3

Figure: Data generated from equilibrium 3

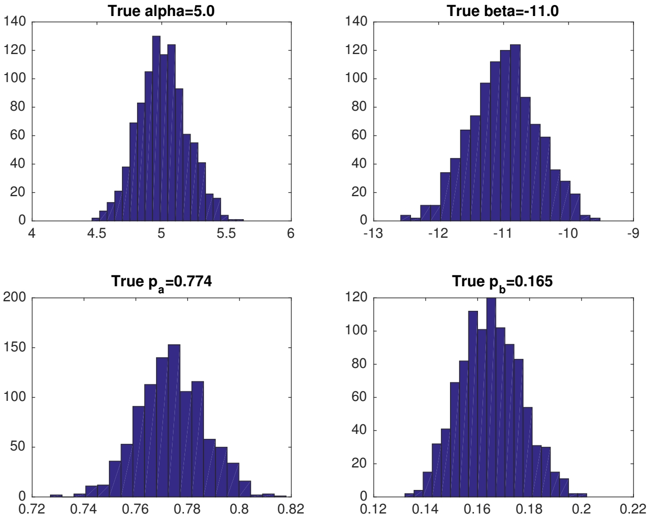

Static Game Example: Maximum Likelihood Estimation

- Maximize the likelihood function

$$
\begin{aligned}
\max_{\alpha,\beta} \quad \log \quad & \mathcal{L}(p_a(\alpha,\beta);X) \\
= \quad & \sum_{i=1}^{N} \left(d_a^i * \log(p_a(\alpha,\beta)) + (1-d_a^i*)\log(1-p_a(\alpha,\beta))\right) \\
+ \quad & \sum_{i=1}^{N} \left(d_b^i * \log(p_b(\alpha,\beta)) + (1-d_b^i*)\log(1-p_b(\alpha,\beta))\right)
\end{aligned}
$$

- $p_a(\alpha,\beta)$ and $p_a(\alpha,\beta)$ are the solutions of the Bayesian-Nash Equilibrium equations

$$
\begin{aligned}
p_a &= \frac{1}{1+\exp[-\alpha x_a + p_b x_a(\alpha-\beta)]} \equiv \Psi_a(p_b, x_a; \alpha, \beta) \\
p_b &= \frac{1}{1+\exp[-\alpha x_b + p_a x_b(\alpha-\beta)]} \equiv \Psi_b(p_a, x_b; \alpha, \beta)
\end{aligned}
$$

Discussion

- Is the likelihood function smooth in α and β for NFXP? What about MPEC - is objective function and constraints smooth in parameters, $\theta = (\alpha, \beta, p_a, p_b)$?
- Sensitivity to starting values?
- Can we identify what equilibrium is played in the data, i.e. the equilibrium selection rule?
- Can we use standard theorems for inference? Is true value in interior of parameter space? Is it differentiable? Is objective function continuous?
- This problem is extremely simple. p_a and p_b are scalars. How would you solve for p_b and p_b when they are solutions to players Bellman equations?
- Can we be sure to find all equilibria by iterating on player's Bellman equations? Why/why not?

Estimation with Multiple Markets

- There 25 different markets, i.e., 25 pairs of observed types $(x_a^m, x_b^m), m = 1, ..., 25$
- The grid on x_a has 5 points equally distributed between the interval $[0.12, 0.87]$, and similarly for x_b
- Use the same true parameter values: (α_0, β_0)
- For each market with (x_a^m, x_b^m), solve BNE conditions for (p_a^m, p_b^m).
- There are multiple equilibria in most of 25 markets
- For each market, we (randomly) choose an equilibrium to generate 1000 data points for that market
- The equilibrium used to generate data can be different in different markets - we flip a coin at each market.

of Equilibria with Different (x_a^m, x_b^m)

Figure: Number of equilibria

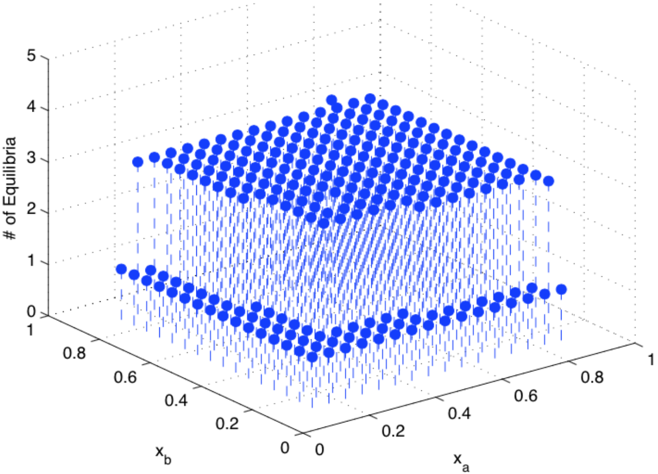

NFXP - Estimation with Multiple Markets

Inner loop:

$$\max_{\alpha,\beta} \quad \log \mathcal{L}(p_a^m(\alpha,\beta), p_b^m(\alpha,\beta); X)$$

Outer loop: For a given values of (α, β) solve BNE equations for ALL equilibria, $k = 1, .., K$ at each market, $m = 1, .., M$: That is, $p_a^{m,k}(\alpha, \beta)$ and $p_a^{m,k}(\alpha, \beta)$ are the solutions to

$$\begin{aligned} p_a^m &= \Psi_a(p_b^m, x_a^m; \alpha, \beta) \\ p_b^m &= \Psi_b(p_a^m, x_b^m; \alpha, \beta) \\ m &= 1, ..., M \end{aligned}$$

where we again choose the equilibrium, that gives the highest likelihood value at each market m

$$k^* = \arg \max_{k=1,\dots,K} \log \mathcal{L}(p_a^{m,k}(\alpha,\beta), p_b^{m,k}(\alpha,\beta); X)$$

such that

$$(p_a^m(\alpha,\beta), p_b^m(\alpha,\beta)) = (p_a^{m,k*}(\alpha,\beta), p_b^{m,k*}(\alpha,\beta))$$

Estimation with Multiple Markets - MPEC

Constrained optimization formulation

$$\max_{\alpha,\beta,p_a^m,p_b^m} \quad \log \mathcal{L}(p_a^m, p_b^m; X)$$

subject to

$$\begin{aligned} p_a^m &= \Psi_a(p_b^m, x_a^m; \alpha, \beta) \\ p_b^m &= \Psi_b(p_a^m, x_b^m; \alpha, \beta) \\ 0 &\leq p_a^m, p_b^m \leq 1, m = 1, ..., M \end{aligned}$$

- MPEC does not explicitly solve the BNE equations to find ALL equilibria at each market - for every trial value of parameters.
- But the number of parameters is much larger.
- Both MPEC and NFXP are based on Full Information Maximum Likelihood (FIML) estimators.

NFXP: Monte Carlo - Multiple Markets (M=25, T=50)

Starting values $\alpha_0 = \alpha$, $\beta_0 = \beta$

Random equilibrium selection in different markets

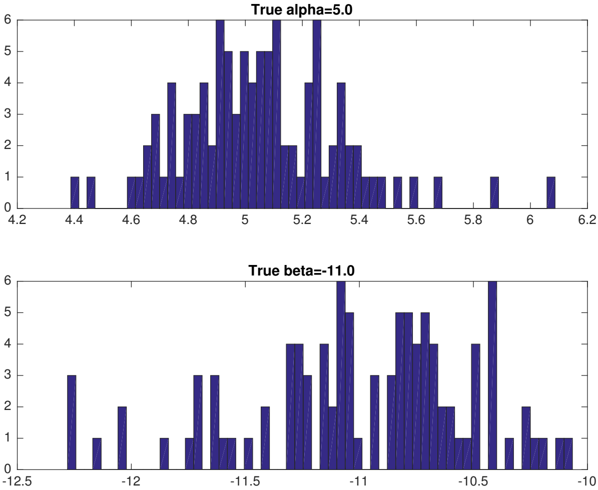

MPEC: Monte Carlo - Multiple Markets (M=25, T=50)

Starting values $\alpha_0 = \alpha$, $\beta_0 = \beta$

Random equilibrium selection in different markets

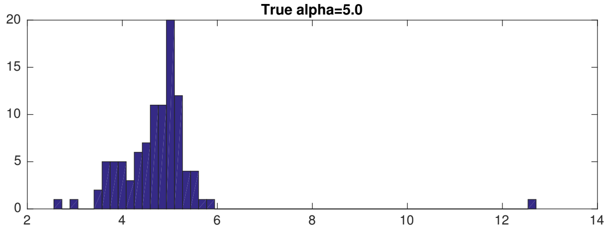

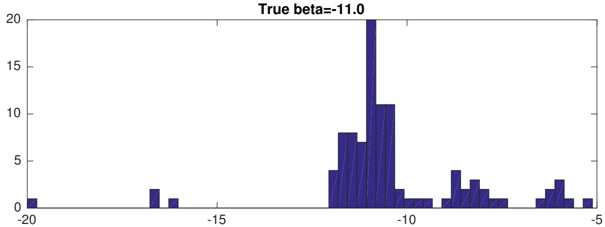

MPEC: Monte Carlo - Multiple Markets (M=2, T=625)

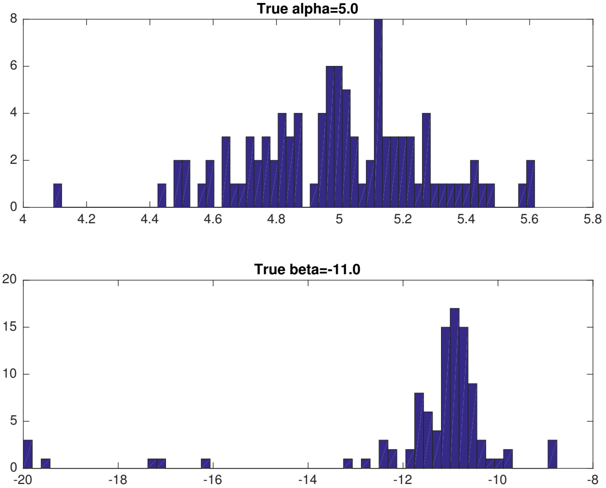

Figure: Random equilibrium selection in different markets

NFXP: Monte Carlo - Multiple Markets (M=25, T=50)

Starting values $\alpha_0 = \alpha$, $\beta_0 = \beta$

Random equilibrium selection in different markets

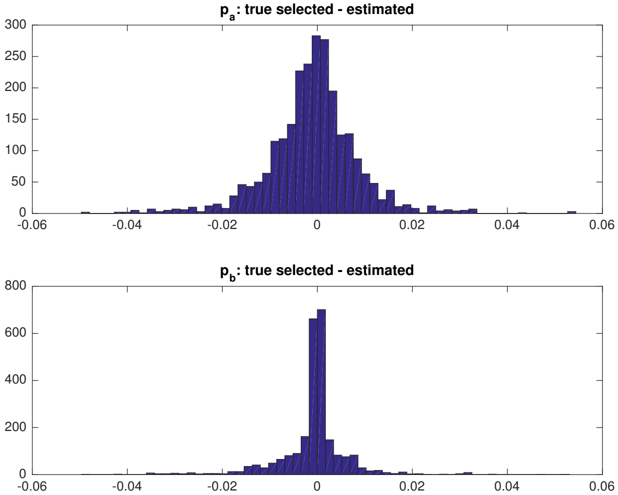

MPEC: Monte Carlo - Multiple Markets (M=25, T=50)

Starting values $\alpha_0 = \alpha$, $\beta_0 = \beta$

Random equilibrium selection in different markets

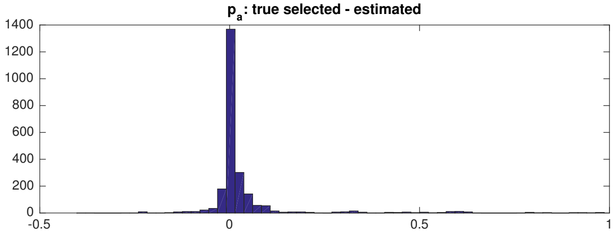

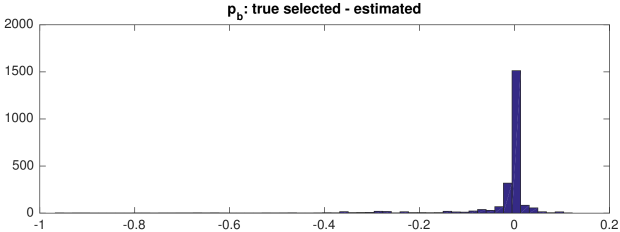

MPEC: Monte Carlo - Multiple Markets (M=2, T=625)

Figure: Random equilibrium selection in different markets

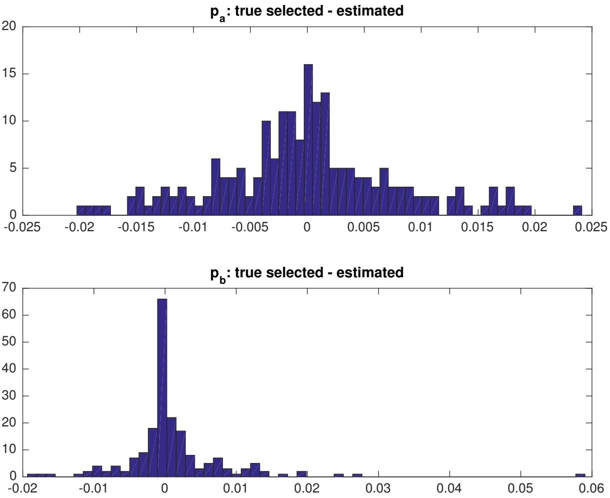

MPEC and NFXP: multiple markets

NFXP:

- 2 parameters in optimization problem
- we can estimate the equilibrium played in the data, $p_a^{m,k*}$ and $p_b^{m,k*}$ (but in models with observationally equivalent equilibria it may not be possible to obtain joint identification of structural parameters and equilibrium probabilities)
- Needs to find ALL equilibria at each market (very hard in more complex problems)
- Good full solution methods required

MPEC:

- $2 + 2M$ parameters in optimization problem
- Does not always converge towards the equilibrium played in the data, although NFXP indicates that $p_a^{m,k*}$ and $p_b^{m,k*}$ are actually identifiable
- Local minima with many markets.
- Disclaimer: Quick and dirty implementation of MPEC. Use AMPL/Knitro

2-Step Methods

Recall the constrained optimization formulation for FIML is

$$\max_{\alpha,\beta,p_a^m,p_b^m} \quad \log \mathcal{L}(p_a^m, p_b^m; X)$$

subject to

$$\begin{aligned} p_a^m &= \Psi_a(p_b^m, x_a^m; \alpha, \beta) \\ p_b^m &= \Psi_b(p_a^m, x_b^m; \alpha, \beta) \\ 0 &\le p_a^m, p_b^m \le 1, m = 1, ..., M \end{aligned}$$

- Denote the solution as $(\alpha^*, \beta^*, p_a^*, p_b^*)$
- Suppose we know (p_a^*, p_b^*), how do we recover (α^*, β^*)?

2-Step Methods: Recovering (α^*, β^*)

- Idea 1: Solve the BNE equations for (α^*, β^*)

$$\begin{aligned} p_a^* &= \Psi_a(p_b^*, x_a; \alpha, \beta) \\ p_b^* &= \Psi_b(p_a^*, x_b; \alpha, \beta) \end{aligned}$$

- Idea 2: Choose (α, β) to

$$\max_{\alpha,\beta} \log \mathcal{L}(\Psi_a(p_b^*, x_a; \alpha, \beta), \Psi_b(p_a^*, x_b; \alpha, \beta); X)$$

2-Step Methods: Recovering (α^*, β^*)

- Idea 1:
  - Step 1: Estimate $\hat{p} = (\hat{p}_a, \hat{p}_b)$ from the data
  - Step 2: Solve

$$\begin{aligned}\hat{p}_a &= \Psi_a(\hat{p}_a, x_a; \alpha, \beta) \\ \hat{p}_b &= \Psi_b(\hat{p}_b, x_b; \alpha, \beta)\end{aligned}$$

- Idea 2
  - Step 1: Estimate $\hat{p} = (\hat{p}_a, \hat{p}_b)$ from the data
  - Step 2: : Choose (α, β) to

$$\max_{\alpha,\beta} \log \mathcal{L}(\Psi_a(\hat{p}_b, x_a; \alpha, \beta), \Psi_b(\hat{p}_a, x_b; \alpha, \beta); X)$$

2-Step Methods: Potential Issues to be Addressed

- How do we estimate $\hat{p} = (\hat{p}_a, \hat{p}_b)$?
- Different methods give different $\hat{p}$
- One method is the frequency estimator:

$$\begin{aligned}\hat{p}_a &= \frac{1}{N}\sum_{i}^{N} I_{\{d_a^i=1\}} \\ \hat{p}_b &= \frac{1}{N}\sum_{i}^{N} I_{\{d_b^i=1\}}\end{aligned}$$

- if $(\hat{p}_a, \hat{p}_b) \neq (p_a^*, p_b^*)$ then $(\hat{\alpha}, \hat{\beta}) \neq (\alpha^*, \beta^*)$
- For a given $(\hat{p}_a, \hat{p}_b)$, there might not be a solution to the BNE equations

$$\begin{aligned}\hat{p}_a &= \Psi_a(\hat{p}_a, x_a; \alpha, \beta) \\ \hat{p}_b &= \Psi_b(\hat{p}_b, x_b; \alpha, \beta)\end{aligned}$$

2-Step Methods: Pseudo Maximum Likelihood

In 2-step methods

- Step 1: Estimate $\hat{p} = (\hat{p}_a, \hat{p}_b)$ from the data
- Step 2: Solve

$$\max_{\alpha,\beta,p_a,p_b} \quad \log \mathcal{L}(p_a, p_b; X)$$

subject to

$$\begin{aligned} p_a &= \Psi_a(\hat{p}_a, x_a; \alpha, \beta) \\ p_b &= \Psi_b(\hat{p}_b, x_b; \alpha, \beta) \\ 0 &\leq p_a^m, p_b^m \leq 1, m = 1, ..., M \end{aligned}$$

Or equivalently

- Step 1: Estimate $\hat{p} = (\hat{p}_a, \hat{p}_b)$ from the data
- Step 2: Solve

$$\max_{\alpha,\beta} \quad \log \mathcal{L}(\Psi_a(\hat{p}_b, x_a; \alpha, \beta), \Psi_b(\hat{p}_a, x_b; \alpha, \beta); X)$$

Least Square Estimators

Pesendofer and Schmidt-Dengler (2008)

- Step 1: Estimate $\hat{p} = (\hat{p}_a, \hat{p}_b)$ from the data
- Step 2: Solve

$$\min_{\alpha,\beta} \left\{ (\hat{p}_a - \Psi_a(\hat{p}_b, x_a; \alpha, \beta))^2 + (\hat{p}_b - \Psi_b(\hat{p}_a, x_b; \alpha, \beta); X))^2 \right\}$$

For dynamic games, Markov perfect equilibrium conditions are characterized by

$$p = \Psi(p, \theta)$$

- Step 1: Estimate $\hat{p}$ from the data
- Step 2: Solve

$$\min_{\alpha,\beta} [\hat{p} - \Psi(\hat{p}; \theta)]' W [\hat{p} - \Psi(\hat{p}; \theta)]'$$

Static Game Example: 2-Step PML

Figure: Data generated from equilibrium 1

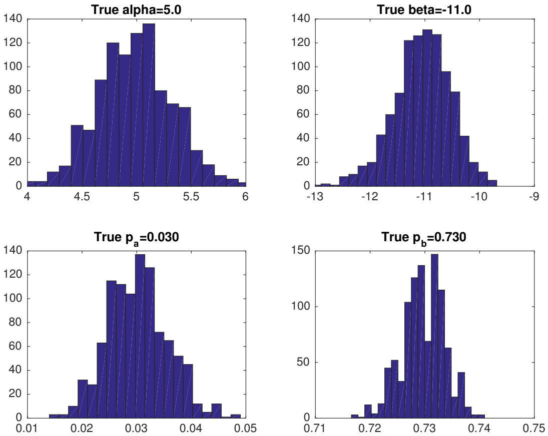

- Pakes, Ostrovsky, and Berry (2007): PML 2-step estimator and can lead to large bias in finite samples.
- In this example, however, it seems to work pretty OK. Why?

Static Game Example: 2-Step PML

Figure: Data generated from equilibrium 2

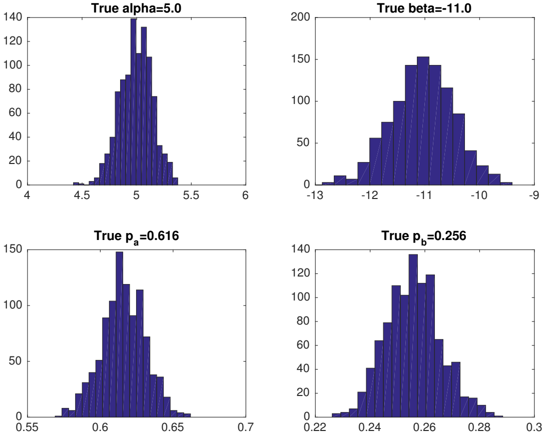

- Pakes, Ostrovsky, and Berry (2007): PML 2-step estimator and can lead to large bias in finite samples.
- In this example, however, it seems to work pretty OK. Why?

Static Game Example: 2-Step PML

Figure: Data generated from equilibrium 3

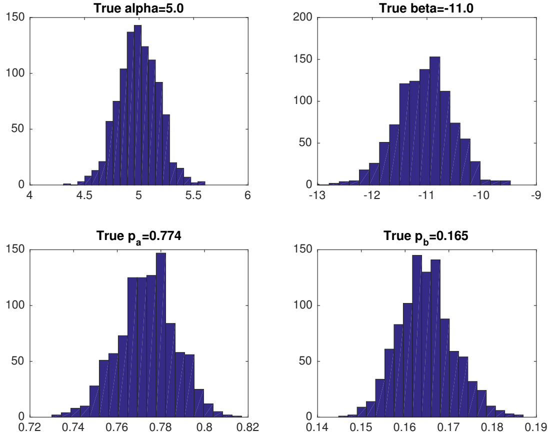

- Pakes, Ostrovsky, and Berry (2007): PML 2-step estimator and can lead to large bias in finite samples.
- In this example, however, it seems to work pretty OK. Why?

Nested Pseudo Likelihood (NPL): Aguirregabiria and Mira (2007)

NPL iterates on the 2-step methods

1. Step 1: Estimate $\hat{p}^0 = (\hat{p}_a^0, \hat{p}_b^0)$ from the data, set $k = 0$
2. Step 2:
   REPEAT
   2.1 Solve
   $$\alpha^{k+1}, \beta^{k+1} = \arg\max_{\alpha,\beta} \quad \log \mathcal{L}(\Psi_a(\hat{p}_b^k, x_a; \alpha, \beta), \Psi_b(\hat{p}_a^k, x_b; \alpha, \beta); X)$$
   2.2 One best-reply iteration on $\hat{p}^k$
   $$\begin{aligned} \hat{p}_a^{k+1} &= \Psi_a(\hat{p}_a^k, x_a; \alpha^{k+1}, \beta^{k+1}) \\ \hat{p}_a^{k+1} &= \Psi_b(\hat{p}_b^k, x_b; \alpha^{k+1}, \beta^{k+1}) \end{aligned}$$
   2.3 Let k:=k+1;
   UNTIL convergence in (α^k, β^k) and $(\hat{p}_a^k, \hat{p}_b^k)$

Monte Carlo Results: NPL with Eq 1

Figure: Equilibrium 1 - $\hat{p}_j = 1/N \sum_i I(d_j = 1)$

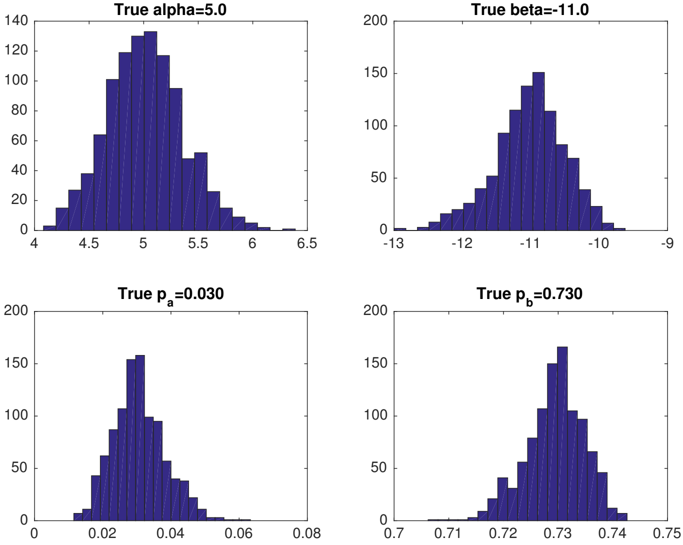

Monte Carlo Results: NPL with Eq 2

Figure: Equilibrium 2 - $\hat{p}_j = 1/N \sum_i I(d_j = 1)$

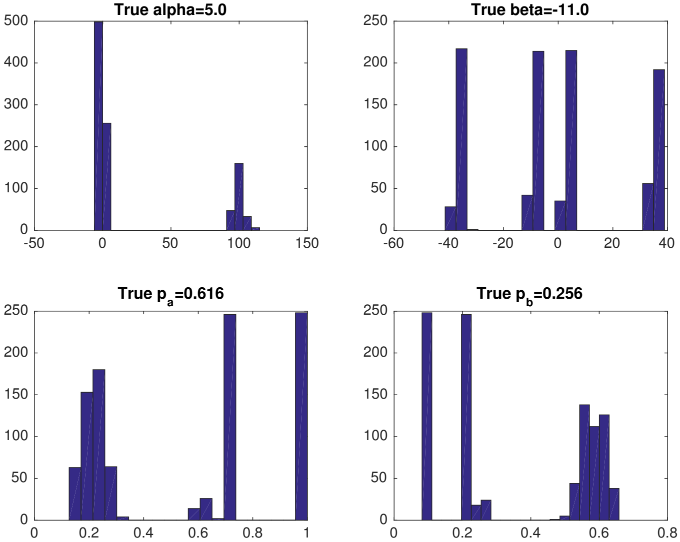

Monte Carlo Results: NPL with Eq 3

Figure: Equilibrium 3 - $\hat{p}_j = 1/N \sum_i I(d_j = 1)$

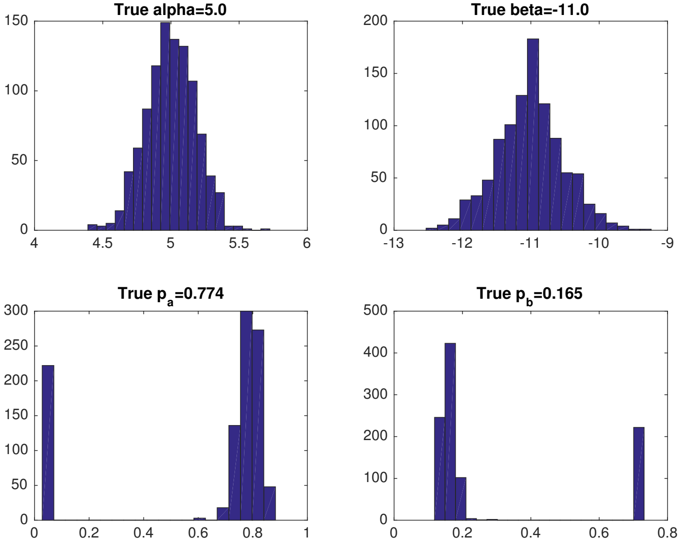

Conclusions

- NFXP/MPEC implementations of MLE is statistically efficient, but computational daunting.
- Two step estimators - computationally fast, but inefficient and biased in small samples.
- NPL (Aguirregabiria and Mira 2007) should bridge this gab, but can be unstable when estimating estimating games with multiple equilibria.
- Estimation of dynamic games is an interesting but challenging computational optimization problem
  - Multiple equilibria leads makes likelihood function discontinuous $\rightarrow$ non-standard inference and computational complexity
  - Multiple equilibria leads to indeterminacy problem and identification issues.
- All these problems are amplified by orders of magnitude when we move to Dynamic models